

Report for computer exercise 4 - Particle simulator in MPI

Anton Karlsson
antka388
931217-7117

1 Code overview

1.1 Discussion on partitioning

The box in which the particles was contained in was partitioned column wise among the processors, meaning that each process owned a local box with the original height and a width of `BOX_VERT_LENGTH/p`, where p is the number of processors. The reasoning behind this partitioning is that it is easy to implement, however, the downside is that the speed-up is not the best. This is because the local boxes that are owned by a process, has an area that decreases fast as the number of processes increase.

A partitioning that might be better is to partition in squares, as the area will not decrease as fast. Thus the probability a particle to a different process will not decrease as fast as the column-wise partitioning. However, the implementation of the square partitioning is harder to code.

1.2 Execution

First, the process defines their dimension of its own local partitioning using the global values of the original box dimensions along with their process id. Then each process initializes equally many particles (given by the parameter `NUM_PART`) with random position (within the local box of the process) and a random velocity. This way, we avoid some initial communication without losing correctness.

The main loop that was given can now be executed. Each process moves its local particles and deal with eventual collisions and wall collisions that occur. Note that any collisions that appear between

particles that are not owned by the same process are ignored. Since the probability of this happening is relatively small and therefore will not affect the final result noticeably.

When the particles have been moved, the process checks whether any particles have coordinates that indicate that they need to move to a neighboring box. The particles that are to be sent to left or right are sent using dynamic communication.

2 Results

Code

Listing 1: main.c

```
#include <stdlib.h>
#include <time.h>
#include <stdio.h>
#include <math.h>
#include <stdbool.h>
#include <mpi.h>

#include "coordinate.h"
#include "definitions.h"
#include "physics.h"

// to do:
// Need to handle nullptr when using plst_t struct.
// Feel free to change this program to facilitate parallelization.
void print_pcord(pcord_t p){

    fprintf(stderr, "x=%f y=%f vx=%f vy=%f\n",
            p.x, p.y, p.vx, p.vy);
}

float rand1(){
    return (float)( rand()/(float) RAND_MAX );
}

void init_collisions(bool *collisions, unsigned int max){
    unsigned int i;
    for(i=0; i<max; ++i)
        collisions[i]=0;
}

bool in_local(pcord_t p, cord_t wall) // checks if particle p is in wall wall.
{
    return (p.x > wall.x0 && p.x < wall.x1);
}

bool exited_left(pcord_t p, cord_t wall, int rank){
    if(rank == 0)
        return false;
    else
        return p.x < wall.x0;
}

bool exited_right(pcord_t p, cord_t wall, int rank, int num_p){
    if(rank == num_p -1)
        return false;
    else
        return p.x > wall.x1;
}

int main(int argc, char** argv){
```

```

double stime = MPI_Wtime();

unsigned int time_stamp = 0, time_max;
int send_left_count=0, send_right_count=0;
int recv_right_count=0, recv_left_count=0;
float pressure = 0, sum = 0;

//MPI stuff
int rank, num_p;
MPI_Init(&argc, &argv);
MPI_Comm comm = MPI_COMM_WORLD;
MPI_Comm_rank(comm, &rank);
MPI_Comm_size(comm, &num_p);
MPI_Status status;
//init_mpi_pcord();

/* create a mpi type for pcord_t */
const int nitems=4;
int blocklengths[4] = {1,1,1,1};
MPI_Datatype ptypes[4] = {MPI_FLOAT, MPI_FLOAT, MPI_FLOAT, MPI_FLOAT};
MPI_Datatype MPI_PCORD;
MPI_Aint poffsets[4];

poffsets[0] = offsetof(pcord_t, x);
poffsets[1] = offsetof(pcord_t, y);
poffsets[2] = offsetof(pcord_t, vx);
poffsets[3] = offsetof(pcord_t, vy);
// poffsets[4] = offsetof(pcord_t, to_remove);

MPI_Type_create_struct(nitems, blocklengths, poffsets, ptypes, &MPI_PCORD);
MPI_Type_commit(&MPI_PCORD);

// Declare send arrays used in mpi functions
//int scounts[num_p];
plst_t *locparticles = NULL;
locparticles = (plst_t *) malloc(sizeof(plst_t));
pcord_t temp_pcord;

// parse arguments
if(argc != 2) {
    fprintf(stderr, "Usage: %s simulation_time\n", argv[0]);
    fprintf(stderr, "For example: %s 10\n", argv[0]);
    exit(-1);
}

time_max = atoi(argv[1]);

/* Initialize */
// 1. set the walls
cord_t wall;
wall.y0 = wall.x0 = 0;
wall.x1 = BOX_HORIZ_SIZE;
wall.y1 = BOX_VERT_SIZE;

cord_t locwall;

```

```

// 2. allocate particle buffer and initialize the particles
bool *collisions= NULL;
collisions = (bool*) malloc(INIT_NO_PARTICLES*num_p*sizeof(bool));

locparticles = NULL;
locparticles = (plst_t*) malloc(INIT_NO_PARTICLES*num_p*sizeof(pcord_t));
pcord_t *send_left = NULL;
send_left = (pcord_t*) malloc(INIT_NO_PARTICLES*num_p*sizeof(pcord_t));
pcord_t *send_right = NULL;
send_right = (pcord_t*) malloc(INIT_NO_PARTICLES*num_p*sizeof(pcord_t));

pcord_t *recv_left = NULL;
recv_left = (pcord_t*) malloc(INIT_NO_PARTICLES*num_p*sizeof(pcord_t));
pcord_t *recv_right = NULL;
recv_right = (pcord_t*) malloc(INIT_NO_PARTICLES*num_p*sizeof(pcord_t));

srand( time(NULL) + 1234 );

// set up walls
locwall.y0 = 0;
locwall.y1 = BOX_VERT_SIZE;
locwall.x0 = rank*BOX_HORIZ_SIZE/num_p;
if (rank == num_p-1)
    locwall.x1 = BOX_HORIZ_SIZE;
else
    locwall.x1 = locwall.x0 + BOX_HORIZ_SIZE/num_p;

float r, a;
int i;

int num_part = INIT_NO_PARTICLES;
// send init particles and designate to processes.
for(i=0; i < num_part; i++) {
    // initialize random position
    temp_pcord.x = locwall.x0 + rand1()*BOX_HORIZ_SIZE/num_p;
    temp_pcord.y = locwall.y0 + rand1()*BOX_VERT_SIZE;

    // initialize random velocity
    r = rand1()*MAX_INITIAL_VELOCITY;
    a = rand1()*2*PI;
    temp_pcord.vx = r*cos(a);
    temp_pcord.vy = r*sin(a);

    push_lst(&locparticles, temp_pcord);
}

/* // print the first elemets */
/* { */
/*   plst_t* current = (plst_t*) malloc(sizeof(plst_t)); */
/*   current = locparticles; */
/*   for (int i = 0; i<5; ++i, current=current->next){ */
/*     fprintf(stderr, "x = %f, y = %f, vx = %f, vy = %f\n", */
/*       current->val.x, current->val.y, current->val.vx, */
/*       current->val.vy); */

```

```

/*    } */
/* } */

// if everything works, this should to

// For each proces, do:
unsigned int p, pp;
int pressure_count =0;

// should reutrn plst_t * (so that a refreence could be returned.)
plst_t* current_part = NULL;
plst_t* sub_current_part = NULL;
current_part = (plst_t*) malloc(sizeof(plst_t*));
sub_current_part = (plst_t*) malloc(sizeof(plst_t*));

/* Main loop */
for (time_stamp=0; time_stamp<time_max; time_stamp++) { // for each time stamp
    init_collisions(collisions, INIT_NO_PARTICLES);
    for(p=0, current_part = locparticles;
        p<num_part;
        p++, current_part = current_part->next) { // for all particles
        //fprintf(stderr, "p = %d for processes with rank %d\n", p, rank);
        //current_part = get_part(locparticles, p);
        if(collisions[p]) continue;
        /* check for collisions */
        // If a local boundary is hit, check if collide with any particle

        for(pp=p+1, sub_current_part=current_part->next;
            pp<num_part;
            pp++, sub_current_part=sub_current_part->next) {
            if(collisions[pp]) continue;

            float t=collide(current_part->val, sub_current_part->val);
            //fprintf(stderr, "t = %1.0f\n", t);
            if(t!=-1){ // collision

                //fprintf(stderr, "collisions in rank %d\n", rank);
                collisions[p]=collisions[pp]=1;
                //fprintf(stderr, "before interact in rank %d\n", rank);
                //print_pcord(current_part->val);
                interact(&(current_part->val), &(sub_current_part->val), t);
                //fprintf(stderr, "after interact in rank %d\n", rank);
                // print_pcord(current_part->val);
                sub_current_part = NULL;
                break; // only check collision of two particles
            }
        }
    }
}

MPI_Barrier(comm);
// move particles that has not collided with another
for(p=send_left_count=send_right_count=0,
    recv_left_count=recv_right_count=0,
    current_part = locparticles;

```

```

    p < num_part;
    ++p, current_part = current_part->next){

if(!collisions[p]) {
    feuler(&(current_part->val), 1);
    pressure += mpi_wall_collide(&(current_part->val),
                                locwall, rank, num_p);

    ++pressure_count;
    //fprintf(stderr, "Pressure was added from rank %d\n", rank);
}
if (exited_left(current_part->val, locwall, rank)) {

    *(send_left + send_left_count) = current_part->val;
    remove_from_list(&locparticles, p);
    --p; --num_part; // to compensate for the removed particle

    ++send_left_count;
    if(num_p == 1 || rank == 0)
        exit(5);
    //printf("sending %d to left\trank %d\n", send_left_count, rank);
}
else if (exited_right(current_part->val, locwall, rank, num_p)) {
    *(send_right + send_right_count) = current_part->val;
    remove_from_list(&locparticles, p);
    --p; --num_part; // to compensate for the removed particle

    ++send_right_count;
    if(num_p == 1 || rank == num_p - 1){
        fprintf(stderr, "The right-most process tried to send a particle to a non-
            existent neighbour to the right\n");
        exit(6);
    }
    //printf("sending %d to right\trank %d\n", send_right_count, rank);
}
}

//Send particles to other boxes

if (rank-1 >= 0) {
    MPI_Send(send_left, send_left_count, MPI_PCORD, rank-1, 0, comm);
    //if (send_left_count > 0)
    //printf("rank %d sent %d particles to left\n", rank, send_left_count);
}

if (rank+1 < num_p) {
    MPI_Send(send_right, send_right_count, MPI_PCORD, rank+1, 1, comm);
    // if (send_right_count > 0)
    // printf("rank %d sent %d particles to right\n", rank, send_right_count);
}

// check if something was sent. Using probe to find recv count.
if(rank-1 >= 0){
    MPI_Probe(rank-1, 1, comm, &status);
    MPI_Get_count(&status, MPI_PCORD, &recv_left_count);
    MPI_Recv(recv_left, recv_left_count, MPI_PCORD, rank-1, 1, comm,

```

```

        &status);
    num_part += recv_left_count;
}

if(rank+1 < num_p){
    MPI_Probe(rank+1, 0, comm, &status);
    MPI_Get_count(&status, MPI_PCORD, &recv_right_count);
    MPI_Recv(recv_right, recv_right_count, MPI_PCORD, rank+1, 0, comm,
             &status);
    num_part += recv_right_count;
}

// update locparticles
// think about what happens when we transfer particles between boxes
multi_push(&locparticles, recv_right, recv_right_count);
multi_push(&locparticles, recv_left, recv_left_count); //prob. need a special case
    hanfle for when recv_right/left == NULL
}

// Gather pressure
int tot_pressure_count;
MPI_Barrier(comm);
MPI_Reduce(&pressure, &sum, 1, MPI_FLOAT, MPI_SUM, 0, comm);
MPI_Reduce(&pressure_count, &tot_pressure_count, 1, MPI_INT, MPI_SUM, 0, comm);

double etime = MPI_Wtime();

if (rank == 0)
    fprintf(stderr, "total time = %g Average pressure = %f pressure count = %d\n",
            etime - stime, sum/(WALL_LENGTH*time_max), tot_pressure_count);

// free stuff
free(collisions);
free(locparticles);
free(send_left);
free(send_right);
free(recv_left);
free(recv_right);
MPI_Type_free(&MPI_PCORD);

MPI_Finalize();
return 0;
}

```